

Road Sign Project Report

Pak Lam Chan, ZI-YING(Sophie) CHEN, Chun Ming Tong

2023-04-25

Contents

Abstract	1
Introduction	2
Background	2
Methodology	3
Data Collection and Annotation:	3
Data Preprocessing:	4
Model Development and Evaluation Metrics:	4
Experiment and Results	6
References	8

Abstract

This report presents the development and evaluation of a machine learning model for road sign detection. The model was trained using the Single Shot Multibox Detector (SSD) architecture and the TensorFlow Object Detection API. A dataset of 100 road sign images, consisting of 5 categories (Stop, Yield, Danger, Deadend, and Parking), was collected and manually labeled. The dataset was split into a training set of 90 images and a testing set of 10 images.

The labeled dataset was preprocessed using Python scripts, and the SSD model was trained using the Stochastic Gradient Descent (SGD) optimizer with a batch size of 4, a learning rate of 0.079999, and a total of 10,000 training steps. The model achieved an average precision (AP) of 0.89 and an average recall (AR) of 0.900 on the testing set.

The model was evaluated using precision and recall metrics on the testing set, and it correctly predicted all road signs in the images except for one, where it misclassified a turn sign as a parking sign. The model's performance can be further improved by collecting more images of the parking sign category.

Overall, the developed machine learning model for road sign detection achieved high accuracy and can be further improved with more data and training. The model has the potential to be integrated into real-world applications for driver assistance and safety.

Introduction

Road signs play an important role in ensuring safety on the roads. The ability to detect and recognize road signs is a crucial aspect of autonomous driving, which has the potential to greatly reduce the number of road accidents caused by human error. In recent years, object detection has become a popular field of research, and many different models have been developed to accurately detect objects in images and videos.

One such model is the Single Shot Detector (SSD), which has been shown to achieve state-of-the-art results in object detection tasks. The SSD model is a type of convolutional neural network (CNN) that is capable of detecting objects at different scales and aspect ratios, making it well-suited for detecting road signs of various sizes and shapes.

In this project, we aim to apply the SSD model to the task of road sign detection. Specifically, we will train the model on a dataset of labeled road sign images and evaluate its performance on a separate test set.

Overall, we believe that the SSD model has great potential for improving the safety of autonomous driving systems by enabling them to accurately detect and recognize road signs. By using advanced computer vision techniques such as object detection, we can take important steps towards a future where roads are safer for everyone.

Background

Road sign recognition is an essential task for safe driving. With the advancement of computer vision technologies, it is possible to automate road sign recognition with the help of object detection models. This project aims to explore the potential of using the Single Shot Multibox Detector (SSD) model to detect road signs in real-time.

Being able to identify road signs with object detection could benefit drivers and road users in different ways. Contemporary vehicles nowadays typically incorporate several cameras installed from different parts of the vehicle body. These peripherals, together with the relevant objection and motion detection solution, enable and empower the delivery of some creative value-added features such as 360-bird-eye-view parking view, lane-deviation and lane-changing warning, and collision warning, among others.

The use of such technologies could make driving safer and more convenient for drivers. Imagine a road condition with low visibility due to weather and night-time or a situation where the driver has demonstrated a certain level of fatigue during driving or is driving on an unfamiliar road. If there is a device or vehicle feature that could detect crucial road signs (e.g., speed limit, rail crossing, construction, etc.) and give appropriate warnings to the drivers, it may potentially reduce traffic accidents and casualties.

While object detection has shown promising results in various applications, it still faces several challenges, especially in road sign detection. One of the primary issues is the diversity of road signs. Road signs can have different shapes, sizes, and colors, and can be partially occluded, rotated, or distorted. Therefore, training a model that can recognize all types of road signs is a challenging task.

Another issue is the real-time performance of the model. Road sign detection should be done in real-time to ensure timely warnings to the driver. Hence, the model should be optimized to provide accurate results while being computationally efficient.

This project is therefore tasked to carry out a preliminary study of potentially useful and applicable models, evaluate and assess the performance and accuracy of the resulting model. Due to the project scope and timeframe, we have focused on a few most commonly selected road signs. We hope that our findings could provide a good starting point and shed some light on the readiness and feasibility of the road sign detection idea, its effectiveness, application constraints and considerations, and its generalization possibilities.

Methodology

Data Collection and Annotation:




Figure 1: Dataset.

Figure 1 illustrates the road sign images dataset we collected on our own from Google Images and manually labeled with the Labeling tool. There are various road signs, and due to our project objective is to make the road safer, we selected 5 categories, *Stop*, *Yield*, *Danger*, *Deadend*, and *Parking* as they represent common road signs that are important for driver safety.

The dataset consisted of 100 images, with each category containing 20 images from different countries with different shapes and graphics, which you can see in Figure 2. We split the dataset into a training set of 90 images and a testing set of 10 images by stratified sampling which ensured that both the training and testing sets had a balanced distribution of images across all categories.



Figure 2: Stop sign from different places. There is white edge for most of the signs, but the one we circled has white line inside.

To create annotations for the dataset, we manually labeled each image by drawing bounding boxes around the road signs in the images and assigned class labels to each box. We then saved the annotations in XML format, with each file corresponding to the annotations for one image. The process is demonstrated in Figure 3.

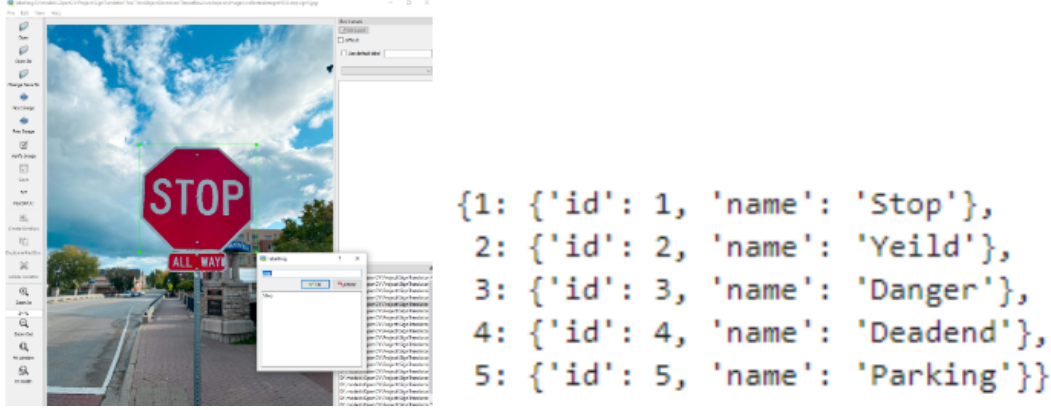


Figure 3: Manually label images.

Data Preprocessing:

To prepare our labeled dataset for training a machine learning model for road sign detection, we used a Python script called `generate_tfrecord.py`. The script takes in several arguments, including the path to the directory containing the XML files, the path to the label map file, the path to the output TFRecord file, and the path to an optional output CSV file.

The script parsed the XML files using the `ElementTree` module in Python, and the resulting data was combined into a Pandas `DataFrame`. We converted the labels in the `DataFrame` from text to integers using a label map dictionary, and split the data into groups based on the filename. Each group was then used to create a TensorFlow example, which was written to the output TFRecord file.

To ensure that all images had the same dimensions and pixel values, we preprocessed the images by resizing them to a fixed size of 300 x 300 pixels and normalizing the pixel values. This step was essential for training our machine learning model, as it ensured that all images had consistent dimensions and pixel values.

Although we did not have the time to augment the images in the training set by applying various transformations such as rotations, flips, and changes in brightness and contrast, we acknowledge that this can help to increase the variety of images that the model sees during training and improve its generalization ability.

Overall, our data processing pipeline successfully converted our annotated dataset into a format suitable for training a machine learning model for road sign detection. By carefully preprocessing the images and converting the labels to integers, we were able to train a model that accurately detected road signs in real-world images.

Model Development and Evaluation Metrics:

We applied Single Shot Detectort model as it is faster and more efficient than traditional models because it only requires a single pass through the image. The model takes an input image with a ground truth bounding box around the target object, and uses anchor boxes in multiple layers to capture the object, illustrates by Figure 4. The model then predicts the category and bounding box offset by comparing the anchor box and true bounding box.

We first developed a base model with batch size 3 and no warm up phase, the result in Figure 5 shows that the average precision and average recall are both 0.427 only, with 5 image categories, low average precision means the model makes many flase positive predictions and low recall means the model makes many false negative predicitons.

After seeing the low values, we decided to increase the batch size which can make the model more rebust because it updates the weights on a more representative samples and incorporated it with a warm-up phase.

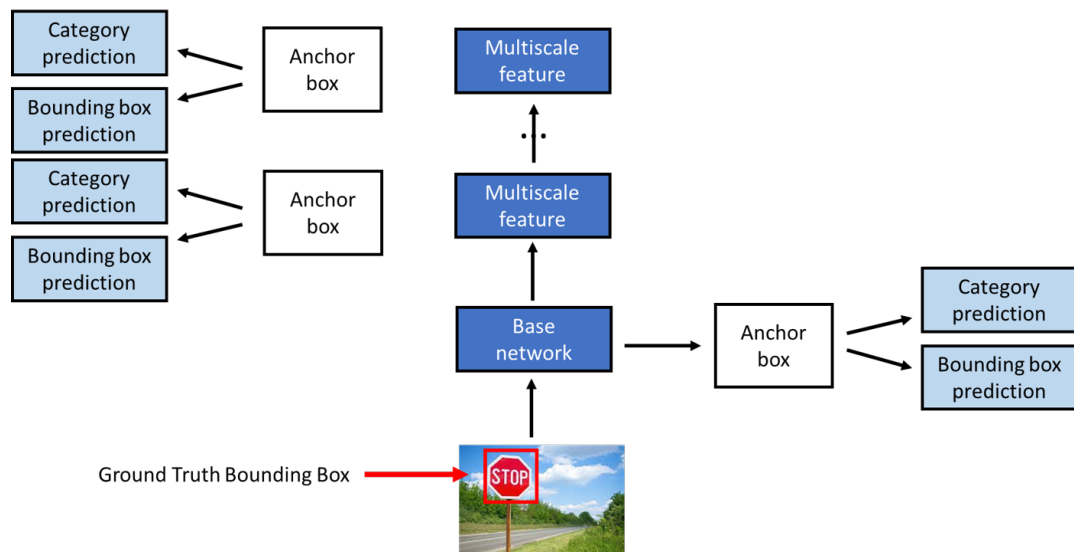


Figure 4: SSD Model

```

DONE (t=0.02s).
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.427
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.625
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.497
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = -1.000
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.427
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.490
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.690
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.690
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = -1.000
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.690
INFO:tensorflow:Eval metrics at step 1000

```

Figure 5: Base Model Training

The final SSD model using the momentum optimizer with a batch size of 4, a learning rate of 0.079999, and a total of 10,000 training steps. We also employed a warm-up stage, where the learning rate was gradually increased from 0 to 0.26666 over the first 1000 training steps to avoid overshooting and local minimum issues.

We used intersection-over-union(IoU) as the measure to detect the object, the threshold was set to be 0.5 to 0.95. Take threshold = 0.5 as an example, when the overlapped area between the predicted bounding box and the true bounding box is higher than 0.5, the object would be detected as the corresponding class of the true bounding box, otherwise, it would be recognized as background object.

We used a mean average precision (mAP) metric to evaluate the model's performance on the test set. The mAP score measures the average precision of the model over multiple levels of recall, which gives a better indication of the model's performance than a single precision or recall score. We calculated the mAP score using the TensorFlow Object Detection API, which computes the score by comparing the predicted bounding boxes to the ground truth boxes.

Overall, we trained the SSD model using SGD with a batch size of 4, a learning rate of 0.079999, and a total of 10,000 training steps, with a warm-up stage of 1000 steps. We fine-tuned the model on our dataset of 90 training images, and evaluated its performance using a mean average precision (mAP) metric on a test set of 10 images.

Experiment and Results

The evaluation metrics for the model performance were based on the five categories of road signs in our dataset, which are Stop, Yield, Danger, Deadend, and Parking. And the average precision and average recall for our base model are both 0.427. In this section, we will only go through the results of our final model and Figure 6 illustrates the model trained after 10000 steps.

Average Precision	(AP) @[IoU=0.50:0.95 area= all maxDets=100]	= 0.890
Average Precision	(AP) @[IoU=0.50 area= all maxDets=100]	= 1.000
Average Precision	(AP) @[IoU=0.75 area= all maxDets=100]	= 1.000
Average Precision	(AP) @[IoU=0.50:0.95 area= small maxDets=100]	= -1.000
Average Precision	(AP) @[IoU=0.50:0.95 area=medium maxDets=100]	= -1.000
Average Precision	(AP) @[IoU=0.50:0.95 area= large maxDets=100]	= 0.890
Average Recall	(AR) @[IoU=0.50:0.95 area= all maxDets= 1]	= 0.900
Average Recall	(AR) @[IoU=0.50:0.95 area= all maxDets= 10]	= 0.900
Average Recall	(AR) @[IoU=0.50:0.95 area= all maxDets=100]	= 0.900
Average Recall	(AR) @[IoU=0.50:0.95 area= small maxDets=100]	= -1.000
Average Recall	(AR) @[IoU=0.50:0.95 area=medium maxDets=100]	= -1.000
Average Recall	(AR) @[IoU=0.50:0.95 area= large maxDets=100]	= 0.900

Figure 6: Model Training

Figure 7 shows the evaluation results on the training set with threshold 0.5 ~ 0.95:

Average Precision (mAP): 0.89

Average Recall: 0.900

Figure 8 shows the evaluation results on the training set with threshold 0.5 ~ 0.95:

Average Precision (mAP): 0.89

Average Recall: 0.900

It is worth noting that the model misclassified one of the testing images, predicting a "Parking" sign instead of a "Turn" sign.(Figure 9) This indicates that the model can be further improved by collecting more images of the "Parking" and "Turn" signs, and potentially adding more training data for those categories.

Overall, the results demonstrate that the SSD model trained on our dataset of road sign images can accurately detect road signs in real-world images with a high degree of precision and recall. The model can be used in various applications, such as driver assistance systems and autonomous vehicles, to improve road safety and reduce the risk of accidents.

Average Precision	(AP) @[IoU=0.50:0.95 area= all maxDets=100]	= 0.890
Average Precision	(AP) @[IoU=0.50 area= all maxDets=100]	= 1.000
Average Precision	(AP) @[IoU=0.75 area= all maxDets=100]	= 1.000
Average Precision	(AP) @[IoU=0.50:0.95 area= small maxDets=100]	= -1.000
Average Precision	(AP) @[IoU=0.50:0.95 area=medium maxDets=100]	= -1.000
Average Precision	(AP) @[IoU=0.50:0.95 area= large maxDets=100]	= 0.890
Average Recall	(AR) @[IoU=0.50:0.95 area= all maxDets= 1]	= 0.900
Average Recall	(AR) @[IoU=0.50:0.95 area= all maxDets= 10]	= 0.900
Average Recall	(AR) @[IoU=0.50:0.95 area= all maxDets=100]	= 0.900
Average Recall	(AR) @[IoU=0.50:0.95 area= small maxDets=100]	= -1.000
Average Recall	(AR) @[IoU=0.50:0.95 area=medium maxDets=100]	= -1.000
Average Recall	(AR) @[IoU=0.50:0.95 area= large maxDets=100]	= 0.900

Figure 7: Model Performance on Training Dataset

Average Precision	(AP) @[IoU=0.50:0.95 area= all maxDets=100]	= 0.890
Average Precision	(AP) @[IoU=0.50 area= all maxDets=100]	= 1.000
Average Precision	(AP) @[IoU=0.75 area= all maxDets=100]	= 1.000
Average Precision	(AP) @[IoU=0.50:0.95 area= small maxDets=100]	= -1.000
Average Precision	(AP) @[IoU=0.50:0.95 area=medium maxDets=100]	= -1.000
Average Precision	(AP) @[IoU=0.50:0.95 area= large maxDets=100]	= 0.890
Average Recall	(AR) @[IoU=0.50:0.95 area= all maxDets= 1]	= 0.900
Average Recall	(AR) @[IoU=0.50:0.95 area= all maxDets= 10]	= 0.900
Average Recall	(AR) @[IoU=0.50:0.95 area= all maxDets=100]	= 0.900
Average Recall	(AR) @[IoU=0.50:0.95 area= small maxDets=100]	= -1.000
Average Recall	(AR) @[IoU=0.50:0.95 area=medium maxDets=100]	= -1.000
Average Recall	(AR) @[IoU=0.50:0.95 area= large maxDets=100]	= 0.900

Figure 8: Model Performance on Testing Dataset



Figure 9: Detection Result of Testing Dataset

References

SSD:

<https://techzizou.com/training-an-ssd-model-for-a-custom-object-using-tensorflow-2-x/>

Tensorflow Object Detection:

https://www.youtube.com/watch?t=10066&v=yqkISICHH-U&feature=youtu.be&ab_channel=NicholasRenotte

The following libraries are adopted in this project:

tensorflow : the major tool used to perform training and inference of deep neural networks. Functionalities also include data automatiom, model tracking, performance monitoring and model training tracking and automation, etc.

google.protobuf : serialize structured data, enable data communication

over networks cv2 : read and write images, alter and translate image properties, filtering amd features detection, etc